Introduction to Linux Session 1 - Basics

Peter Ruprecht peter.ruprecht@colorado.edu

www.rc.colorado.edu

Outline

- What is Linux?
- Why use Linux?
- What happens when you log in?
- Shells and environment
- Commands
- Filesystem basics
- Processes
- More about shells

Slides available at

https://github.com/ResearchComputing/ USGS_2014-07

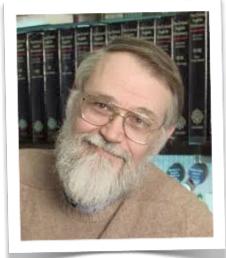
What is Linux?

- Part of the Unix family of operating systems.
- Started in early '90s by Linus Torvalds.
- Technically refers only to the kernel; software from the GNU project and elsewhere is layered on top to form a complete OS. Most is open source.
- Several distributions are available from commercialgrade, like RHEL or SUSE, to more consumer-focused, like Ubuntu.
- Runs on everything from embedded systems to supercomputers.

Why Use Linux?

- Linux command-line syntax may seem overwhelming to the new user, but:
- It's the default operating system on virtually all HPC systems
- It's extremely flexible
- It tries not to get in your way
- It's fast and powerful
- It was designed by programmers and thus has many potent tools for software development
- You can get started with a few basic commands and build from there

History of Linux



Brian Kernighan 1970 "space travel" to Unix



Dennis Ritchie 1971



Richard Stallman 1983 Gnu Not Unix



Linus Torvalds 1991 Linux kernel for personal computers

users

shell: bash, csh programs commands Linux kernel Computer hardware

How do you log in?

- To a remote system, use Secure Shell (SSH)
- From Windows GUI app such as PuTTY
- From Linux ssh on the command line
 ssh –X username@login.rc.colorado.edu
- From Mac OS X ssh from the Terminal, or GUI such as Cyberduck or Fugu

What happens when you log in?

- Login is authenticated (password or key)
- Assigned to a tty
- Shell starts
- Environment is set up
- Prompt

What identifies a Linux user?

- Username / UUID
- Group / GID
- Password (or other authentication info)
- GECOS
- Default shell
- Home directory

Shells

The shell parses and interprets typed input; passes results to the rest of the OS; returns response as appropriate

- Bourne (sh) early and rudimentary
- Bourne-again (bash) has many user-friendly extensions; default in Linux
- C (csh) has C-like syntax
- T (tcsh) extended version of C
- Korn (ksh) early extension of Bourne; was heavily used for programming
- Z (zsh) includes features of bash and tcsh

Shell features

- Tab completion
- History and command-line editing
- Scripting and programming
- Built-in utilities

Environment

- Set up using shell and environment variables
 - shell: only effective in the current shell itself
 - environment: carry forward to subsequent commands or shells
- Set default values at login time using .bash_profile (or .profile). Non-login interactive shells will read .bashrc instead. Use .my.bashrc in RC!
- Initialization scripts should not produce output!
- set var_name[=value] (shell)
- export VAR_NAME[=value] (environment)
- env (shows current variables)

Useful variables

- PATH: directories to search for commands
- HOME: home directory
- DISPLAY: screen where graphical output will appear
- MANPATH: directories to search for manual pages
- LANG: current language encoding
- PWD: current working directory
- USER: username
- LD_LIBRARY_PATH: directories to search for shared objects (dynamically-loaded libs)
- LM_LICENSE_FILE: files to search for FlexLM software licenses

Anatomy of a Linux command

- Command [flags] [flag arguments] [target(s)]
- tar —c —f archive.tar mydir
- Flags do not mean the same thing for different commands
- The same command may have different flags in different kinds of Unix (esp. Linux vs BSD)
- Case is important!
- Order of flags may be important

Most important Linux command

man

man <command>
man -k <keyword>

File- and directory-related commands

```
pwd – prints full path to current directory
cd – changes directory; can use full or relative path as target
mkdir – creates a subdirectory in the current directory
rmdir – removes an empty directory
rm - removes a file (rm -r removes a directory and all of its
contents)
cp – copies a file
mv – moves (or renames) a file
Is – lists the contents of a directory (1s –1 gives detailed
listing)
chmod/chown – change permissions or ownership
df – displays filesystems and their sizes
du – shows disk usage (du –sk shows size of a directory and
all of its contents in KB)
```

Process- and program-related commands

ps – lists processes (ps –ef lists all running processes)

top – shows processes currently using the CPU

kill – sends a signal to a process (kills process by default). Target is Process-ID; found in 2nd column of ps —ef

output.

jobs - shows jobs currently in background

time – shows how much wall time and CPU time a process has used

nice - changes the priority of a process to get CPU time

File-viewing commands

less – displays a file one screen at a time

cat – prints entire file to the screen

head – prints the first few lines of a file

tail – prints the last few lines of a file (with –f shows in realtime the end of a file that may be changing)

diff – shows differences between two files

grep – prints lines containing a string or other regular expression

tee – prints the output of a command and also copies the output to a file

sort – sorts lines in a file

find – searches for files that meet specified criteria

wc - count words, lines, or characters in a file

The Linux Filesystem

- System of arranging files on disk
- Consists of directories (folders) that can contain files or other directories
- Levels in full paths separated by forward slashes, e.g. /home/admin/mary/payroll/June2012
- Case-sensitive; spaces in names discouraged
- ., .., and ~ are shorthand.

Much more on this in the next session!

Navigating the filesystem

- Examples:
 - Is
 - mkdir
 - cd
 - rm
- Permissions (modes)

File editing

- nano simple and intuitive to get started with; not powerful; keyboard driven
- vi/vim universal; keyboard driven; powerful but some learning curve required
- emacs keyboard or GUI versions; helpful extensions for programmers; well-documented
- OpenOffice / LibreOffice for WYSIWYG

http://xkcd.com/378/

Processes

- A process is a unique task; it may have threads
- Examples:
 - Foreground vs background (&)
 - jobs command
 - Ctl-C vs Ctl-Z ; bg
 - kill

More about shells

- Input and output redirection
 - Send output from a command to a new file with >
 - Append output to an existing file with >>
 - Use a file as input to a command with <
- Pipes: | sends output of one command to another command

```
ps —ef | grep ruprech
```

Quoting – save this for a future session!

Thank you!

Slides available at:

http://researchcomputing.github.io